

```
In [1]: install.packages('leaps')
```

package 'leaps' successfully unpacked and MD5 sums checked

The downloaded binary packages are in

C:\Users\ashka\AppData\Local\Temp\RtmpkFu7iv\downloaded_packages

```
In [2]: library(caret)
library(leaps)
library(MASS)
library(glmnet)
```

Loading required package: lattice

Loading required package: ggplot2

Registered S3 methods overwritten by 'ggplot2':

| method | from |
|----------------|-------|
| [.quosures | rlang |
| c.quosures | rlang |
| print.quosures | rlang |

Loading required package: Matrix

Loading required package: foreach

Loaded glmnet 2.0-16

```
In [3]: raw_data<-read.table(file = './uscrime.txt',header = TRUE)
```

```
In [4]: #It is necessary to scale the data as mentioned also in the lecture to optimize the coefficient range
```

```
remove_points <- c("So", "Crime")
data_remove_points<-raw_data[, !(names(raw_data) %in% remove_points)]
scaled_data_removed<-scale(data_remove_points)
scaled_data<-cbind(scaled_data_removed,raw_data[,remove_points])
```

```
In [5]: # The step wise regression is implemented here:
# The range of features is tried to find the optimum model

control_variable <- trainControl(method = "repeatedcv", number = 5, repeats = 5)

step_wise_model <- train(Crime ~., data = scaled_data, method = "leapSeq", tuneGrid = data.frame(nvmax = 1:15),
trControl = control_variable)
```

```
In [6]: step_wise_model$results
```

| nvmax | RMSE | Rsquared | MAE | RMSESD | RsquaredSD | MAESD |
|-------|----------|-----------|----------|----------|------------|----------|
| 1 | 291.5779 | 0.5227764 | 230.8915 | 80.52834 | 0.2889516 | 74.80276 |
| 2 | 308.1104 | 0.4503600 | 238.4573 | 93.82773 | 0.3121311 | 71.89600 |
| 3 | 265.5925 | 0.5740250 | 209.2237 | 77.24358 | 0.2716775 | 68.41373 |
| 4 | 285.1170 | 0.5081315 | 221.7260 | 61.58699 | 0.2441184 | 48.16886 |
| 5 | 276.3819 | 0.5224353 | 222.6152 | 77.18560 | 0.2469174 | 62.99584 |
| 6 | 246.7716 | 0.5991393 | 193.8343 | 60.97895 | 0.2226581 | 50.10355 |
| 7 | 272.6054 | 0.5435213 | 220.7917 | 59.79278 | 0.2280076 | 49.79663 |
| 8 | 280.2471 | 0.5340670 | 229.3941 | 61.54290 | 0.2449292 | 51.32562 |
| 9 | 301.7002 | 0.4500372 | 240.6683 | 65.08583 | 0.2498655 | 57.50237 |
| 10 | 286.1816 | 0.5202469 | 236.0677 | 54.98126 | 0.2345198 | 48.36711 |
| 11 | 289.8436 | 0.5142951 | 234.0024 | 52.93549 | 0.2307848 | 45.59564 |
| 12 | 290.2612 | 0.5073331 | 234.4481 | 54.11271 | 0.2037706 | 44.13606 |
| 13 | 281.5523 | 0.5313749 | 226.9553 | 49.46382 | 0.2164079 | 41.29760 |
| 14 | 279.3022 | 0.5386282 | 225.1971 | 47.90591 | 0.2026085 | 36.52641 |
| 15 | 291.6086 | 0.5022123 | 235.3427 | 47.57263 | 0.2122186 | 37.83192 |

In [7]: `summary(step_wise_model)`

Subset selection object
15 Variables (and intercept)

| | Forced in | Forced out |
|--------|-----------|------------|
| M | FALSE | FALSE |
| Ed | FALSE | FALSE |
| Po1 | FALSE | FALSE |
| Po2 | FALSE | FALSE |
| LF | FALSE | FALSE |
| M.F | FALSE | FALSE |
| Pop | FALSE | FALSE |
| NW | FALSE | FALSE |
| U1 | FALSE | FALSE |
| U2 | FALSE | FALSE |
| Wealth | FALSE | FALSE |
| Ineq | FALSE | FALSE |
| Prob | FALSE | FALSE |
| Time | FALSE | FALSE |
| So | FALSE | FALSE |

1 subsets of each size up to 6

Selection Algorithm: 'sequential replacement'

| | M | Ed | Po1 | Po2 | LF | M.F | Pop | NW | U1 | U2 | Wealth | Ineq | Prob | Time | So |
|---------|---|----|-----|-----|----|-----|-----|----|----|----|--------|------|------|------|----|
| 1 (1) | " | " | " | " | " | " | " | " | " | " | " | " | " | " | " |
| 2 (1) | " | " | " | " | " | " | " | " | " | " | " | " | " | " | " |
| 3 (1) | " | " | " | " | " | " | " | " | " | " | " | " | " | " | " |
| 4 (1) | " | " | " | " | " | " | " | " | " | " | " | " | " | " | " |
| 5 (1) | " | " | " | " | " | " | " | " | " | " | " | " | " | " | " |
| 6 (1) | " | " | " | " | " | " | " | " | " | " | " | " | " | " | " |

In [8]: *#After investigating which features are selected by the step_wise*
#It is possible to built a linear regression based on the selected factors
`linear_regression_model<-lm(formula = Crime ~ M+So+Ed + Po1 , data = scaled_data)`

In [9]: linear_regression_model

Call:

```
lm(formula = Crime ~ M + So + Ed + Po1, data = scaled_data)
```

Coefficients:

| (Intercept) | M | So | Ed | Po1 |
|-------------|--------|--------|-------|--------|
| 855.90 | 132.78 | 144.49 | 92.05 | 314.42 |

Lasso Regression

In [10]: *# The following developed data set will be used for the Lasso Regression*
 head(scaled_data)

| | M | Ed | Po1 | Po2 | LF | M.F | Pop | NW | U1 | U2 | Wealth |
|--|------------|------------|------------|------------|------------|-------------|-------------|--------------|-------------|------------|------------|
| | 0.9886930 | -1.3085099 | -0.9085105 | -0.8666988 | -1.2667456 | -1.12060499 | -0.09500679 | 1.943738564 | 0.69510600 | 0.8313680 | -1.3616094 |
| | 0.3521372 | 0.6580587 | 0.6056737 | 0.5280852 | 0.5396568 | 0.98341752 | -0.62033844 | 0.008483424 | 0.02950365 | 0.2393332 | 0.3276683 |
| | 0.2725678 | -1.4872888 | -1.3459415 | -1.2958632 | -0.6976051 | -0.47582390 | -0.48900552 | 1.146296747 | -0.08143007 | -0.1158877 | -2.1492481 |
| | -0.2048491 | 1.3731746 | 2.1535064 | 2.1732150 | 0.3911854 | 0.37257228 | 3.16204944 | -0.205464381 | 0.36230482 | 0.5945541 | 1.5298536 |
| | 0.1929983 | 1.3731746 | 0.8075649 | 0.7426673 | 0.7376187 | 0.06714965 | -0.48900552 | -0.691709391 | -0.24783066 | -1.6551781 | 0.5453053 |
| | -1.3983912 | 0.3898903 | 1.1104017 | 1.2433590 | -0.3511718 | -0.64550313 | -0.30513945 | -0.555560788 | -0.63609870 | -0.5895155 | 1.6956723 |

In [11]: *# The Lasso model will be developed in the following:*

```
lasso_model = cv.glmnet(x = as.matrix(scaled_data[, -16]),
                        y = as.matrix(scaled_data$Crime),
                        alpha = 1, nfolds = 5,
                        type.measure = "mse", family = "gaussian", standardize = F)
```

In [12]: *#The summary of the lasso model is presented in the following:*

```
coef(lasso_model, s = lasso_model$lambda.min)
```

16 x 1 sparse Matrix of class "dgCMatrix"

| | 1 |
|-------------|-----------|
| (Intercept) | 905.08511 |
| M | 87.38594 |
| Ed | 123.07830 |
| Po1 | 308.82944 |
| Po2 | . |
| LF | . |
| M.F | 52.10363 |
| Pop | . |
| NW | 11.13820 |
| U1 | -28.85703 |
| U2 | 62.03082 |
| Wealth | . |
| Ineq | 187.18594 |
| Prob | -77.09109 |
| Time | . |
| So | . |

In [13]: *# Regression model for the lasso will be presented in the following:*

```
lasso_linear_model <- lm(Crime ~ M + Ed + Po1 + M.F + NW + U1 + U2 + Ineq + Prob, data = scaled_data)
```

```
In [14]: # Testing the trained lasso model after developing a linear regression
summary(lasso_linear_model)
```

Call:

```
lm(formula = Crime ~ M + Ed + Po1 + M.F + NW + U1 + U2 + Ineq +
    Prob, data = scaled_data)
```

Residuals:

| | Min | 1Q | Median | 3Q | Max |
|--|--------|--------|--------|-------|-------|
| | -439.2 | -102.2 | -6.3 | 124.1 | 476.6 |

Coefficients:

| | Estimate | Std. Error | t value | Pr(> t) | |
|-------------|----------|------------|---------|----------|-----|
| (Intercept) | 905.09 | 28.87 | 31.352 | < 2e-16 | *** |
| M | 111.23 | 46.83 | 2.375 | 0.022820 | * |
| Ed | 203.63 | 60.12 | 3.387 | 0.001687 | ** |
| Po1 | 297.89 | 52.08 | 5.719 | 1.51e-06 | *** |
| M.F | 68.74 | 41.63 | 1.651 | 0.107134 | |
| NW | 16.55 | 53.15 | 0.311 | 0.757222 | |
| U1 | -109.46 | 60.94 | -1.796 | 0.080609 | . |
| U2 | 156.94 | 62.09 | 2.528 | 0.015889 | * |
| Ineq | 236.70 | 61.95 | 3.821 | 0.000492 | *** |
| Prob | -89.99 | 36.28 | -2.481 | 0.017791 | * |

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 197.9 on 37 degrees of freedom

Multiple R-squared: 0.7894, Adjusted R-squared: 0.7381

F-statistic: 15.41 on 9 and 37 DF, p-value: 4.881e-10

Elastic Net

```
In [17]: #It is important to find the alpha value to be able to develop the Elastic net

mse_list <- numeric()

search_alpha <- function(num_value, scaled_data){
  alpha <- num_value
  elastic_net <- cv.glmnet(x=as.matrix(scaled_data[,-16]),
                          y=as.matrix(scaled_data[,16]),
                          alpha = alpha,
                          nfolds=5,
                          type.measure="mse",
                          family="gaussian",
                          standardize=FALSE)

  mse_list <-< cbind(mse_list, c(alpha, min(elastic_net$cvm),elastic_net$lambda.min))
}
```

```
In [18]: for (i in seq(.01,1,by = .01)){search_alpha(i,scaled_data)}
```

```
In [19]: minIndex <- which.min(mse_list[2,])
```

```
In [20]: # The alpha value will be used as part of the model for the development.
```

```
mse_list[1, minIndex]
```

0.8

```
In [21]: elastic_net_final <- cv.glmnet(x=as.matrix(scaled_data[,-16]),
    y=as.matrix(scaled_data[,16]),
    alpha = 0.26,
    nfolds=5,
    type.measure="mse",
    family="gaussian",
    standardize=FALSE)
```

```
In [22]: coef(elastic_net_final, s = elastic_net_final$lambda.min)

# The linear regression model will be develop in the following section
elastic_linear_regression <- lm(Crime ~ ., data = scaled_data[, -4])
```

16 x 1 sparse Matrix of class "dgCMatrix"

| | 1 |
|-------------|------------|
| (Intercept) | 900.219667 |
| M | 93.139696 |
| Ed | 138.787373 |
| Po1 | 199.220825 |
| Po2 | 78.485074 |
| LF | . |
| M.F | 66.718527 |
| Pop | -2.658487 |
| NW | 27.340856 |
| U1 | -64.139269 |
| U2 | 99.200515 |
| Wealth | 39.484214 |
| Ineq | 200.575065 |
| Prob | -85.446760 |
| Time | . |
| So | 14.292230 |

In [23]: *# The summary of the elastic linear regression will be presented in the following:*

```
summary(elastic_linear_regression)
```

Call:

```
lm(formula = Crime ~ ., data = scaled_data[, -4])
```

Residuals:

| | Min | 1Q | Median | 3Q | Max |
|--|---------|---------|--------|--------|--------|
| | -442.55 | -116.46 | 8.86 | 118.26 | 473.49 |

Coefficients:

| | Estimate | Std. Error | t value | Pr(> t) | |
|-------------|----------|------------|---------|----------|-----|
| (Intercept) | 903.155 | 58.889 | 15.336 | 2.66e-16 | *** |
| M | 112.934 | 52.244 | 2.162 | 0.038232 | * |
| Ed | 198.350 | 68.044 | 2.915 | 0.006445 | ** |
| Po1 | 286.864 | 71.091 | 4.035 | 0.000317 | *** |
| LF | -11.321 | 56.896 | -0.199 | 0.843538 | |
| M.F | 53.684 | 59.798 | 0.898 | 0.376026 | |
| Pop | -29.833 | 48.950 | -0.609 | 0.546523 | |
| NW | 25.149 | 63.619 | 0.395 | 0.695239 | |
| U1 | -97.649 | 75.332 | -1.296 | 0.204164 | |
| U2 | 143.034 | 69.378 | 2.062 | 0.047441 | * |
| Wealth | 87.540 | 99.662 | 0.878 | 0.386292 | |
| Ineq | 290.076 | 90.023 | 3.222 | 0.002921 | ** |
| Prob | -97.432 | 49.655 | -1.962 | 0.058484 | . |
| Time | -7.991 | 47.425 | -0.168 | 0.867251 | |
| So | 5.669 | 148.100 | 0.038 | 0.969705 | |

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 208.6 on 32 degrees of freedom

Multiple R-squared: 0.7976, Adjusted R-squared: 0.709

F-statistic: 9.006 on 14 and 32 DF, p-value: 1.673e-07

In []: